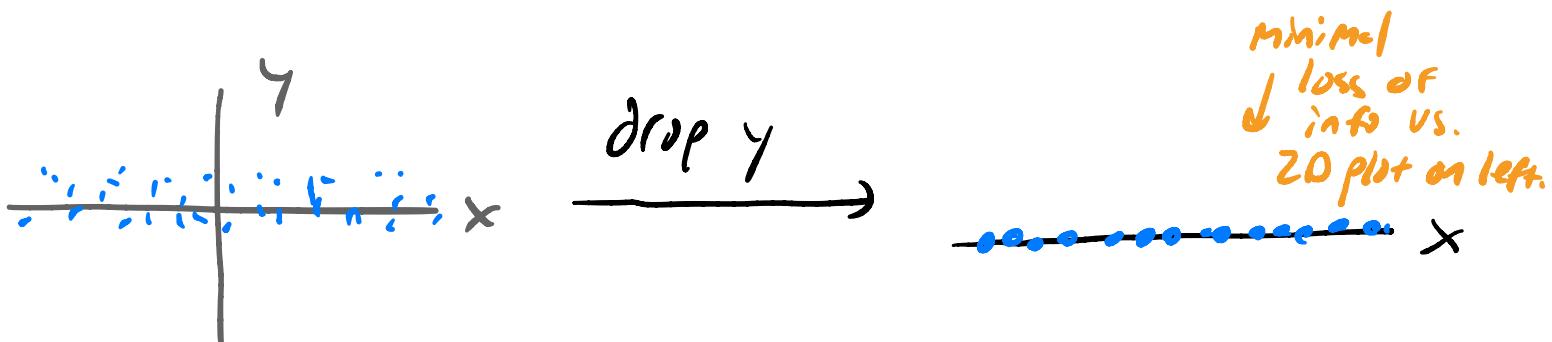
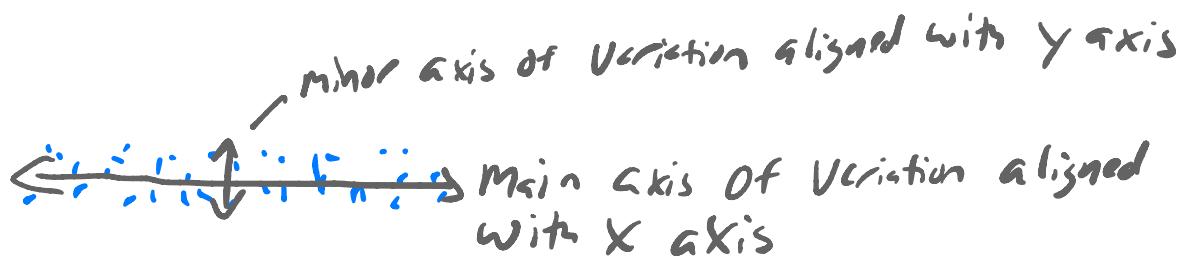


Principal Component Analysis (PCA)

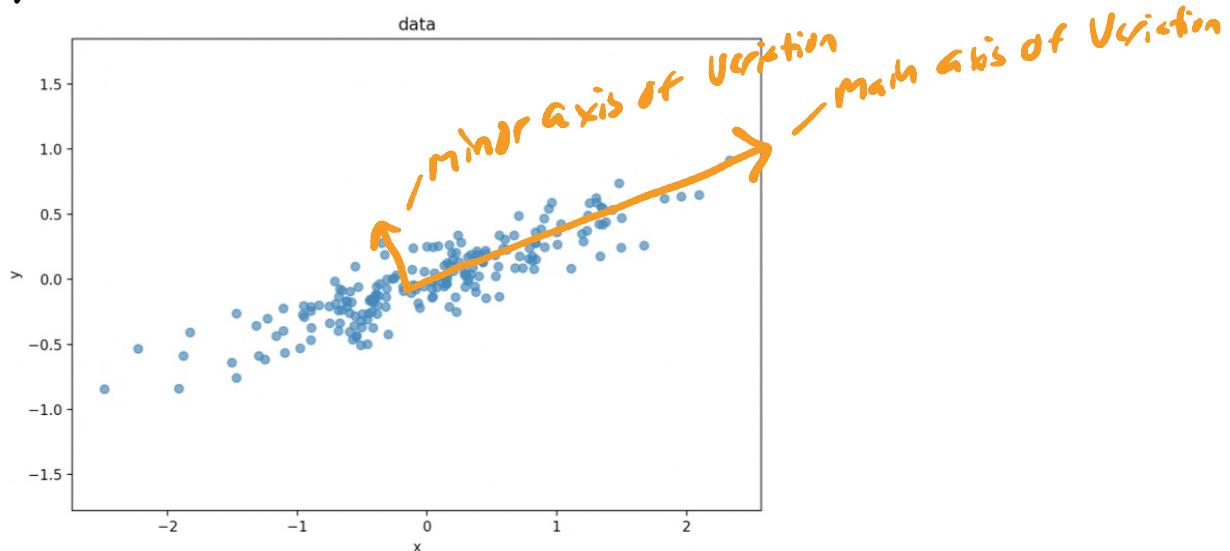
Recall the example where we could toss out a Variable (y) without lossing much information:



In this case, the main axes of variation of data are aligned with the coordinate axes



\Rightarrow clearly this is not true in general! Axes of variation do not need to correspond to variables in dataset
— e.g. $x, y, \text{etc.}$:

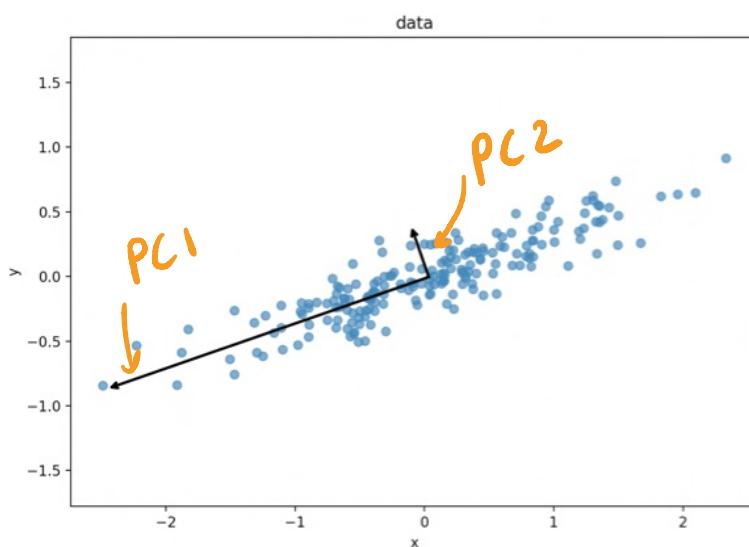


Above Main Variation Axis != x axis!
minor Variation axis != y axis!

We call these intrinsic axes of variation in data
principal components (PCs)

⇒ We number them in order based on the amount
of variation along each intrinsic axis

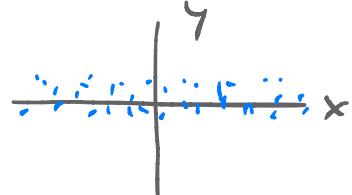
⇒ Larger variation → Smaller number.



Goal of PCA

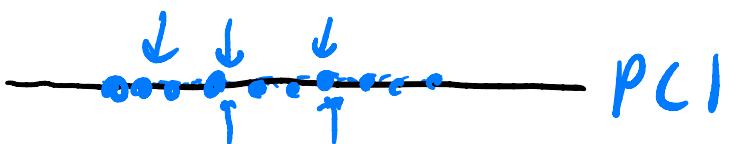
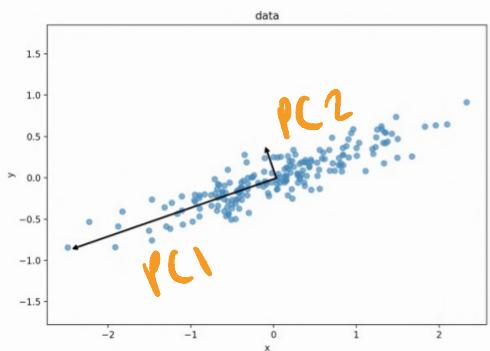
Reduce dimensionality of dataset
(e.g. $M=500 \rightarrow 30$) by dropping
intrinsic axes with little variation
(least amount of info lost)

e.g. PC2 above, y variable in



As we do not want to drop data variables

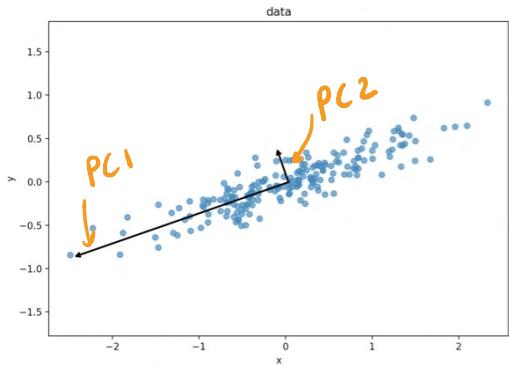
e.g. example drop PC 2.



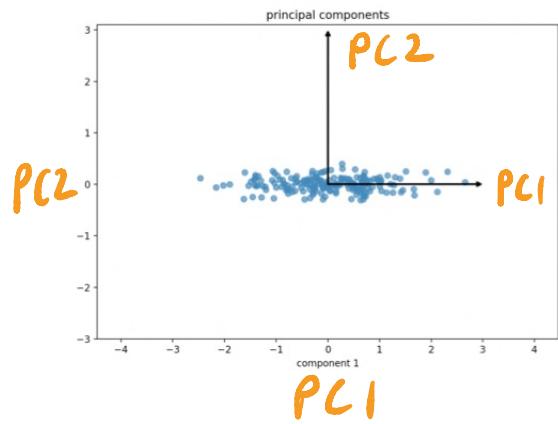
How do we "access" the PCs? — they can be oriented in any direction (but they are assumed to be orthogonal to each other).

⇒ Rotate the data so that the PCs become coordinate axes:

New coordinate axes; The axes of primary variation in the data



Rotate
→



How is this done? A translation then a rotation matrix applied:

See plots above with centered data at $(0,0)$

[Rotated, centered data]

rotation matrix

centered data matrix: $A_c = A - \bar{\mu}$

translation

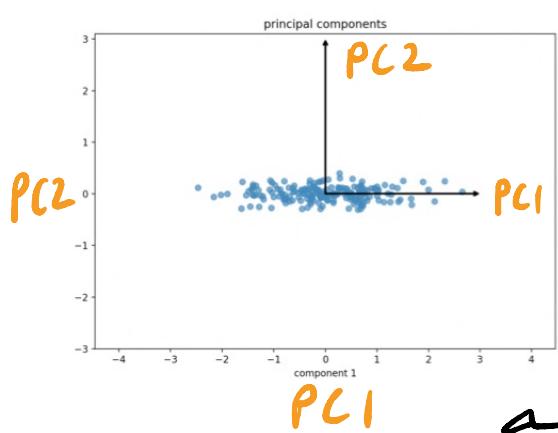
$\hat{A}_c = (P \cdot T @ A_c) \cdot T$

or more concisely:

$\hat{A}_c = A_c @ P$

Same thing

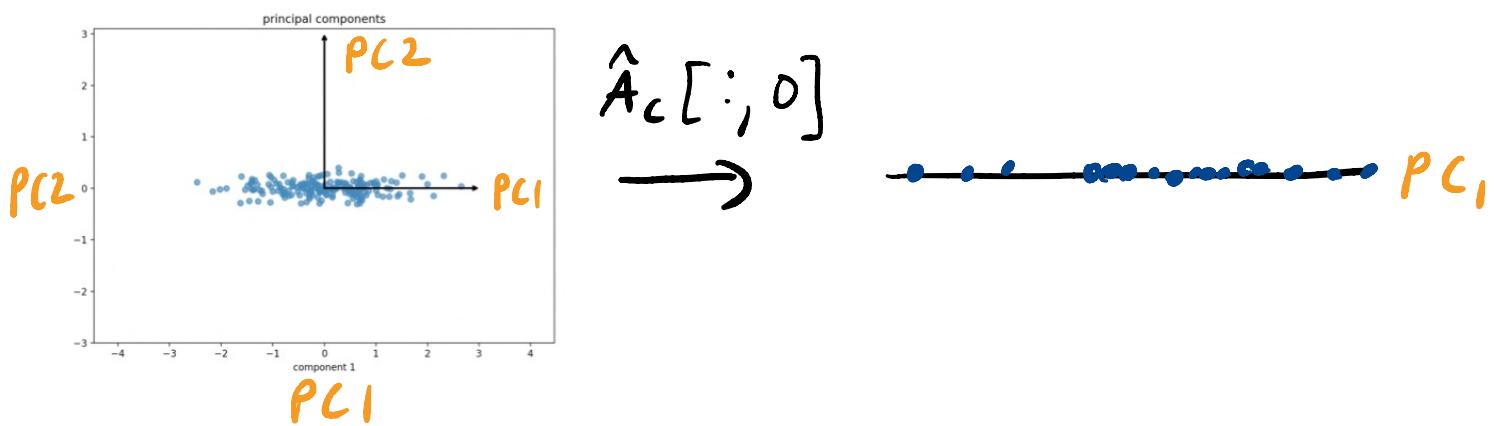
\hat{A}_c



\hat{A}_c is centered data in PCA Space: coordinate axes are PC1 & PC2, not X, Y

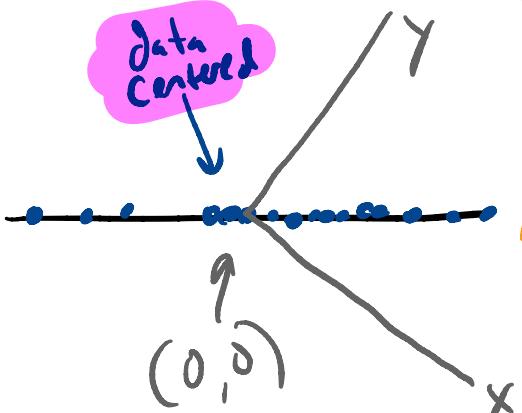
... and to drop the insignificant PCs, we can project onto the large meaningful ones (e.g. PC1)

Remember: project = index vers — e.g. $\hat{A}_c[:, 0]$



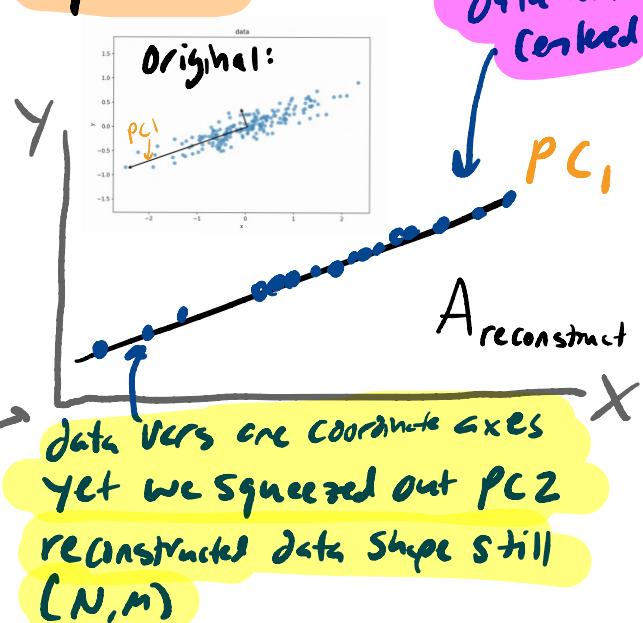
This achieves what we wanted, but the coordinate axes are PCs — not data variables, which might not be what we want.

We can express this result in terms of the original data variables (reconstruct the data) by undoing the rotation + translation!

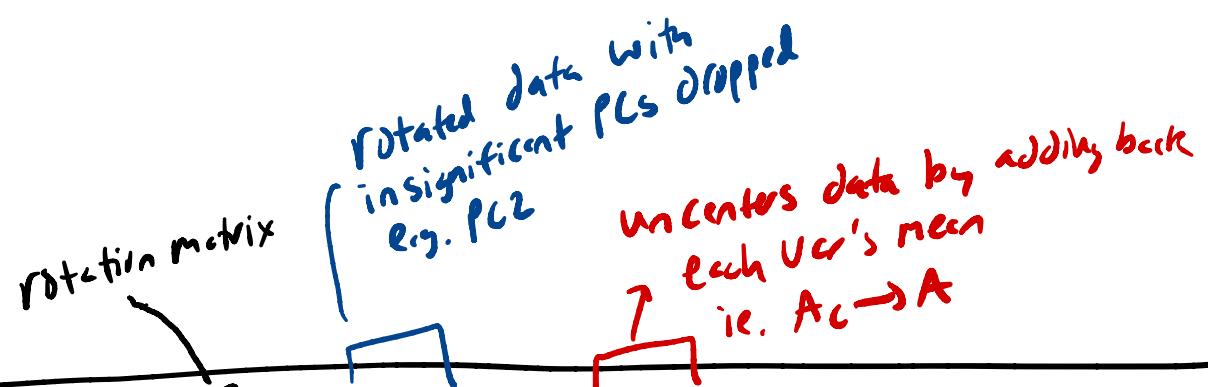


reverse
R + T

(0,0)



To reverse R & T and reconstruct data from \hat{A}_c :



$$A_{\text{reconstruct}} = (\hat{P} @ \hat{A}_c . T) . T + \vec{\mu} = \hat{A}_c @ \hat{P} . T + \vec{\mu}$$

Same — use either equation

Q: What is the rotation matrix P ?

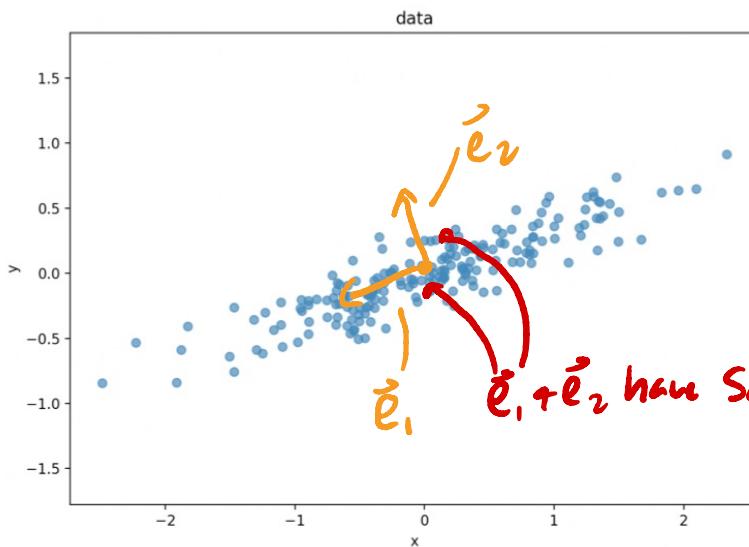
P is the eigenvectors of the Covariance matrix Σ of the data.

$$\text{Recall: } \Sigma = \frac{1}{N-1} A_c . T @ A_c$$

Shape: (n, n)

M of them — one per data variable

Eigen Vectors $P = [\vec{e}_1, \vec{e}_2, \dots, \vec{e}_n]$
 tell us the direction but not amount of the
 corresponding PCs.



* Rotating data by P aligns coordinate axes with PCs

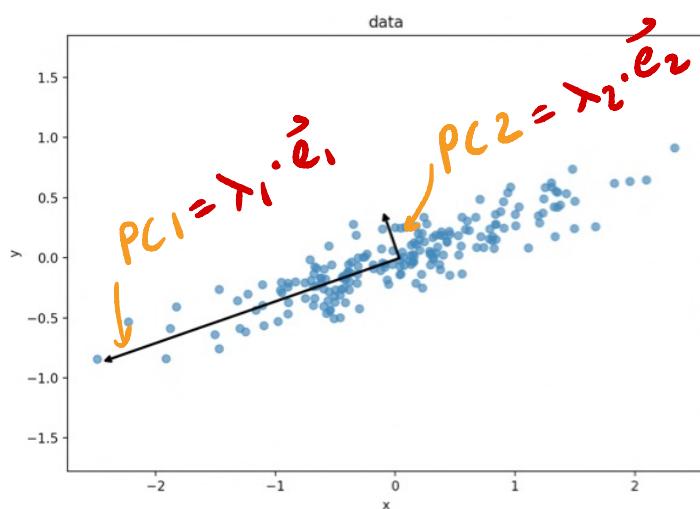
How do we know the amount of Variation in each PC

Direction? Eigen Values = $[\lambda_1, \lambda_2, \dots, \lambda_n]$

M Scalars/floats

One per PC/Eigen vector/Data Variable

$\leftarrow \lambda_1$ is much bigger than λ_2 !



IV Computing eigenvectors and eigenvalues

In Numpy:

$$e_vals, e_vecs = \text{np. linalg. eig}(\Sigma)$$

Computes the **Spectral decomposition** of Σ

P is an orthogonal matrix (forms orthogonal basis)

V Strategy for removing insignificant PCs:

e_vals tell us amount of variation in a PC
 \Rightarrow if small, we can remove.

e.g. $e_vals = [0.1, 3, 0.7, 5]$

\Rightarrow Sort high \rightarrow low:

$$e_vals = [5, 3, 0.7, 0.1]$$

These account for small part of overall data variation, might be good candidate to remove PC 3 & PC 4 here.

★ How many PCs to drop / what is significant is up to us to decide.

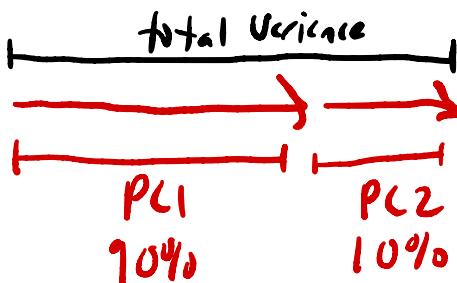
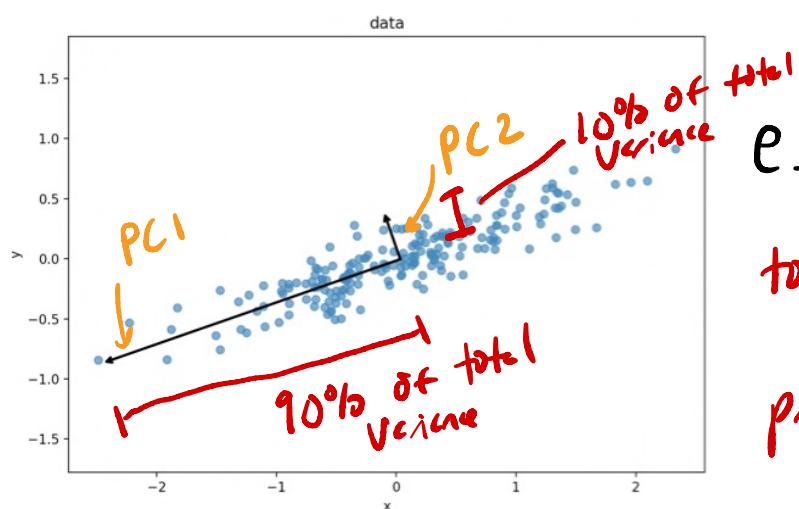
⇒ useful to determine the

proportion of total Variance Accounted for by each PC

$$e\text{-Vals} = [5, 3, 0.7, 0.1]$$

(expressing each as a fraction of the total Variance "pot" more helpful:

these numbers have little meaning on their own



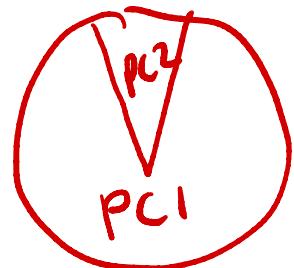
★ prop-Var must sum to 100% (or 1).

$$e\text{-Vals} = [5, 3, 0.7, 0.1]$$

$$\begin{aligned} \text{total Var} &= 5 + 3 + 0.7 + 0.1 \\ &= 8.8 \end{aligned}$$

$$\begin{aligned} \text{prop-Var} &= \left[\underbrace{\frac{5}{8.8}}, \underbrace{\frac{3}{8.8}}, \underbrace{\frac{0.7}{8.8}}, \underbrace{\frac{0.1}{8.8}} \right] \\ &\quad \text{57\%} \quad \text{34\%} \quad \text{8\%} \quad \text{1\%} \end{aligned}$$

$$\text{Sum} = 100\%$$



Often helpful to define a **Variance - kept Cut-off (threshold)** — e.g. toss out PCs to keep at least 90% of the total Variance in the data.

⇒ Compute cumulative Variance accounted for :

Cum-Vcr =

$$\begin{aligned} & \left[\frac{5}{8.8}, \frac{3}{8.8} + \frac{5}{8.8}, \frac{0.7}{8.8} + \frac{3}{8.8} + \frac{5}{8.8}, \frac{0.1}{8.8} + \frac{0.7}{8.8} + \frac{3}{8.8} + \frac{5}{8.8} \right] \\ &= \left[0.57, 0.9, \underbrace{0.99}_{\text{Keep}}, \underbrace{1.0}_{\text{Discard}} \right] \end{aligned}$$

after accounting for all PCs including this smallest one, total Vcr must be 1.0 i.e. 100%.

To keep 90% of the information in data / Variance, only keep top 2 PCs.

Summary of PCA

- 1) [optional] Normalize data if appropriate so that each data value in range $[0, 1]$.

Note: Can be done in vectorized way without matrix multiplication with one line of code.

- **Normalize per-Variable (separately)** may be appropriate if data variables have many different units (e.g. miles & cms)
⇒ You lose relative scales among vars
- **normalize globally (together)** may be appropriate if variables have similar units (e.g. inches)
⇒ Won't distort relative scales among vars

- 2) Center data by subtracting each variable by its mean $\vec{\mu}$

$$A_c = A - \vec{\mu}$$

- 3) Compute covariance matrix: $\Sigma = \frac{1}{N-1} A_c \cdot T \in A_c$

- 4) Compute eigenvectors P and eigenvalues via spectral decomposition of Σ — via Numpy.

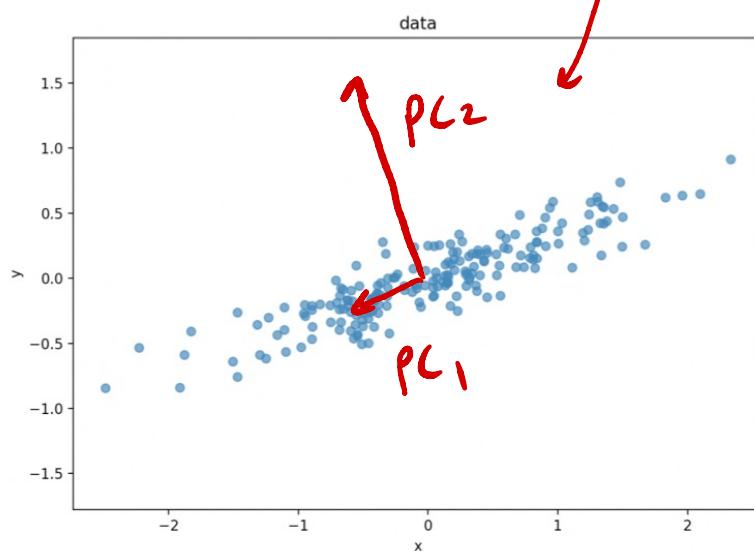
Shape: (M, M)

5) Sort e-vals high \rightarrow low

e.g. $[3.2, 1.1, 0.03]$

Sort P **(columns)** in the same order otherwise PC lengths & directions get Scrambled!

problem if we don't
sort P + evals
the same way!



6a) Compute proportion Variance accounted for by each PC

e.g. total = $3.2 + 1.1 + 0.03 = 4.33$

$\Rightarrow [3.2/4.33, 1.1/4.33, 0.03/4.33]$

6b) Compute cumulative Variance accounted for by PC_i :

e.g. $[3.2/4.33, (3.2+1.1)/4.33, (3.2+1.1+0.03)/4.33]$

6c) Threshold PCs based on desired amount of info/variance to keep. We keep K PCs.

e.g. $\geq 90\%$. If $PC1 + PC2 = 90\%$, toss out $PC3$.

7) project data into PCA Space [rotating data to align coordinate axes to PCs] \hat{P} is P with certain PCs / cols dropped

$$\hat{A}_c = \underbrace{\hat{A}_c}_{(N, K)} @ \underbrace{\hat{P}}_{(N, M)} \quad \text{keep top } k \text{ PCs}$$

8) [optional] Reconstruct data [rotating data after projecting away insignificant PCs to align coordinate axes with original data variables]

if you didn't normalize:

$$\underbrace{A_{\text{reconstruct}}}_{(N, M)} = \underbrace{\hat{A}_c}_{(N, K)} @ \underbrace{\hat{P}^T}_{(K, M)} + \vec{\mu}$$

if you normalized:

$$A_{\text{reconstruct}} = \vec{s} \cdot (\hat{A}_c @ \hat{P}^T) + \vec{\mu}$$

Original data range in each variable